

This slide was intentionally left blank.

# Why You Might Want To Go Async

*PyCon Indonesia 2017 - Jonas Obrist - HDE Inc*

Hi, I'm Jonas

Hi, I'm Jonas



Hi, I'm Jonas



# Hi, I'm Jonas



# Quick Links

[@ojiidotch](#)

[github.com/ojii](https://github.com/ojii)

[hde.co.jp/en/](https://hde.co.jp/en/)

# Why You Might Want To Go Async



~~Why You Might Want To Go Async~~

~~Why You Might Want To Go Async~~

What Do I Mean When I Say Async

*Async*

*Asynchronous*

# Asynchronous IO

# Asynchronous Networking

Asynchronous Networking  
Opposite of Synchronous

# Examples of Synchronous (Python) Frameworks/Libraries



# Examples of Synchronous (Python) Frameworks/Libraries

- `requests`

# Examples of Synchronous (Python) Frameworks/Libraries

- `requests`
- `Django`

# Examples of Synchronous (Python) Frameworks/Libraries

- `requests`
- `Django`
- `Flask`

# Examples of Synchronous (Python) Frameworks/Libraries

- `requests`
- `Django`
- `Flask`
- `psycopg2` (PostgreSQL Library)

# Examples of Synchronous (Python) Frameworks/Libraries

- `requests`
- `Django`
- `Flask`
- `psycopg2` (PostgreSQL Library)
- `boto3` (AWS Library)

# Examples of Synchronous (Python) Frameworks/Libraries

- `requests`
- `Django`
- `Flask`
- `psycopg2` (PostgreSQL Library)
- `boto3` (AWS Library)
- `<most other libraries here>`

# Examples of Asynchronous (Python) Frameworks/Libraries

# Examples of Asynchronous (Python) Frameworks/Libraries

- `twisted`



# Examples of Asynchronous (Python) Frameworks/Libraries

- `twisted`
- `tornado`

# Examples of Asynchronous (Python) Frameworks/Libraries

- `twisted`
- `tornado`
- `asyncio` (stdlib)

# Examples of Asynchronous (Python) Frameworks/Libraries

- `twisted`
- `tornado`
- `asyncio` (stdlib)
- `sanic` (“Flask”)

# Examples of Asynchronous (Python) Frameworks/Libraries

- `twisted`
- `tornado`
- `asyncio` (stdlib)
- `sanic` (“Flask”)
- `curio`

# Examples of Asynchronous (Python) Frameworks/Libraries

- `twisted`
- `tornado`
- `asyncio` (stdlib)
- `sanic` (“Flask”)
- `curio`
- `aiopg` (PostgreSQL Library)

# Examples of Asynchronous (Python) Frameworks/Libraries

- `twisted`
- `tornado`
- `asyncio` (stdlib)
- `sanic` (“Flask”)
- `curio`
- `aiopg` (PostgreSQL Library)
- `aiobotocore` (AWS Library)

# Examples of Asynchronous (Python) Frameworks/Libraries

- `twisted`
- `tornado`
- `asyncio` (stdlib)
- `sanic` (“Flask”)
- `curio`
- `aiopg` (PostgreSQL Library)
- `aiobotocore` (AWS Library)
- `arsenic` & `aapns` (Libraries I wrote and am shamelessly plugging here)

# Asynchronous IO Core Concepts



# Asynchronous IO Core Concepts

- Event Loop

# Asynchronous IO Core Concepts

- Event Loop
- Coroutines (Previously: Callback Hell)

# Asynchronous IO Core Concepts

- Event Loop
- Coroutines (Previously: Callback Hell)
  - A Function which allows cooperative multi-tasking (Co-operative Sub-Routine)

# Asynchronous IO Core Concepts

- Event Loop
- Coroutines (Previously: Callback Hell)
  - A Function which allows cooperative multi-tasking (Co-operative Sub-Routine)
- IO is non-blocking

# Asynchronous IO Core Concepts

- Event Loop
- Coroutines (Previously: Callback Hell)
  - A Function which allows cooperative multi-tasking (Co-operative Sub-Routine)
- IO is non-blocking
- “Logic” might be blocking

# Why You Might Want To Go Async



# Why You Might Want To Go Async



~~Why You Might Want To Go Async~~

~~Why You Might Want To Go Async~~

Why Does Going Async Save Money

# Python Is “Slow”

# Python Is “Slow”

- Synchronous web server can only handle **one** request per thread/process.

# Python Is “Slow”

- Synchronous web server can only handle **one** request per thread/process.
- Scaling is usually done with more threads/processes/servers

# Python Is “Slow”

- Synchronous web server can only handle **one** request per thread/process.
- Scaling is usually done with more threads/processes/servers
- Lots of resources (CPU) is wasted

# Python Is “Slow”

- Synchronous web server can only handle **one** request per thread/process.
- Scaling is usually done with more threads/processes/servers
- Lots of resources (CPU) is wasted
  - Most of the time we're waiting for something to respond with data

# Python Is “Slow”

- Synchronous web server can only handle **one** request per thread/process.
- Scaling is usually done with more threads/processes/servers
- Lots of resources (CPU) is wasted
  - Most of the time we're waiting for something to respond with data
  - Surprisingly little time is spent in our actual Python code



Python Is “Slow”

~~Python Is “Slow”~~

~~Python Is “Slow”~~

Python Is “Inefficient”

~~Python Is “Slow”~~

~~Python Is “Inefficient”~~

~~Python Is “Slow”~~

~~Python Is “Inefficient”~~

Synchronous Python Is Inefficient

*Asynchronous Python To The Rescue!*

# Example: A Web Server

# Example: A Web Server

- Handle a `POST` request



# Example: A Web Server

- Handle a `POST` request
- Store the `POST` data in a database

# Example: A Web Server

- Handle a `POST` request
- Store the `POST` data in a database
- Return an identifier of the database object as json

# Example: Webserver (Synchronous)

# Example: Webserver (Synchronous)

```
def handler(request: Request) -> Response:
    data = request.post_data
    pk = database.save_item(data)
    return Response(
        status=200,
        body=json.dumps({'pk': pk}) ,
        content_type='application/json')
)
```

# Example: Webserver (Synchronous)

```
def handler(request: Request) -> Response: # First, Second
    data = request.post_data
    pk = database.save_item(data)
    return Response(
        status=200,
        body=json.dumps({'pk': pk}) ,
        content_type='application/json')
)
```

# Example: Webserver (Synchronous)

```
def handler(request: Request) -> Response: # Second
    data = request.post_data                # First
    pk = database.save_item(data)
    return Response(
        status=200,
        body=json.dumps({'pk': pk}) ,
        content_type='application/json')
)
```

# Example: Webserver (Synchronous)

```
def handler(request: Request) -> Response: # Second
    data = request.post_data
    pk = database.save_item(data)             # First
    return Response(
        status=200,
        body=json.dumps({'pk': pk}) ,
        content_type='application/json')
)
```

# Example: Webserver (Synchronous)

```
def handler(request: Request) -> Response: # Second
    data = request.post_data
    pk = database.save_item(data)
    return Response(                          # First
        status=200,
        body=json.dumps({'pk': pk}) ,
        content_type='application/json')
)
```



# Example: Webserver (Synchronous)

```
def handler(request: Request) -> Response:
    data = request.post_data
    pk = database.save_item(data)
    return Response(
        status=200,
        body=json.dumps({'pk': pk}) ,
        content_type='application/json')
    )
# First Done
# Second
```

# Example: Webserver (Synchronous)

```
def handler(request: Request) -> Response:
    data = request.post_data
    pk = database.save_item(data)                # Second
    return Response(
        status=200,
        body=json.dumps({'pk': pk}) ,
        content_type='application/json')
    )
# First Done
```

# Example: Webserver (Synchronous)

```
def handler(request: Request) -> Response:
    data = request.post_data
    pk = database.save_item(data)
    return Response(                                     # Second
        status=200,
        body=json.dumps({'pk': pk}) ,
        content_type='application/json')
    )
# First Done
```

# Example: Webserver (Synchronous)

```
def handler(request: Request) -> Response:
    data = request.post_data
    pk = database.save_item(data)
    return Response(
        status=200,
        body=json.dumps({'pk': pk}) ,
        content_type='application/json')
    )
# First Done
# Second Done
```

# Example: Webserver (Asynchronous)

# Example: Webserver (Asynchronous)

```
async def handler(request: Request) -> None:
    data = await request.get_post_data()
    pk = await database.save_item(data)
    response = request.response(
        status=200,
        content_type='application/json',
    )
    response.write(json.dumps({'pk': pk}))
    response.finish()
```

# Example: Webserver (Asynchronous)

```
async def handler(request: Request) -> None:
    data = await request.get_post_data()
    pk = await database.save_item(data)
    response = request.response(
        status=200,
        content_type='application/json',
    )
    response.write(json.dumps({'pk': pk}))
    response.finish()
```

# Example: Webserver (Asynchronous)

```
async def handler(request: Request) -> None:
    data = await request.get_post_data()
    pk = await database.save_item(data)
    response = request.response(
        status=200,
        content_type='application/json',
    )
    response.write(json.dumps({'pk': pk}))
    response.finish()
```



# Example: Webserver (Asynchronous)

```
async def handler(request: Request) -> None:
    data = await request.get_post_data()
    pk = await database.save_item(data)
    response = request.response(
        status=200,
        content_type='application/json',
    )
    response.write(json.dumps({'pk': pk}))
    response.finish()
```

# Example: Webserver (Asynchronous)

```
async def handler(request: Request) -> None:
    data = await request.get_post_data()
    pk = await database.save_item(data)
    response = request.response(
        status=200,
        content_type='application/json',
    )
    response.write(json.dumps({'pk': pk}))
    response.finish()
```

# Example: Webserver (Asynchronous)

```
async def handler(request: Request) -> None: # First, Second
    data = await request.get_post_data()
    pk = await database.save_item(data)
    response = request.response(
        status=200,
        content_type='application/json',
    )
    response.write(json.dumps({'pk': pk}))
    response.finish()
```

# Example: Webserver (Asynchronous)

```
async def handler(request: Request) -> None:    # Second
    data = await request.get_post_data()         # First
    pk = await database.save_item(data)
    response = request.response(
        status=200,
        content_type='application/json',
    )
    response.write(json.dumps({'pk': pk}))
    response.finish()
```

# Example: Webserver (Asynchronous)

```
async def handler(request: Request) -> None:
    data = await request.get_post_data()
    pk = await database.save_item(data)
    response = request.response(
        status=200,
        content_type='application/json',
    )
    response.write(json.dumps({'pk': pk}))
    response.finish()
```

# First, Second

# Example: Webserver (Asynchronous)

```
async def handler(request: Request) -> None:
    data = await request.get_post_data()           # First
    pk = await database.save_item(data)            # Second
    response = request.response(
        status=200,
        content_type='application/json',
    )
    response.write(json.dumps({'pk': pk}))
    response.finish()
```

# Example: Webserver (Asynchronous)

```
async def handler(request: Request) -> None:
    data = await request.get_post_data()
    pk = await database.save_item(data)
    response = request.response(
        status=200,
        content_type='application/json',
    )
    response.write(json.dumps({'pk': pk}))
    response.finish()
```

# First  
# Second

# Example: Webserver (Asynchronous)

```
async def handler(request: Request) -> None:
    data = await request.get_post_data()
    pk = await database.save_item(data)
    response = request.response(                                # First
        status=200,
        content_type='application/json',
    )
    response.write(json.dumps({'pk': pk}))
    response.finish()
    # Second Done
```



# Example: Webserver (Asynchronous)

```
async def handler(request: Request) -> None:
    data = await request.get_post_data()
    pk = await database.save_item(data)
    response = request.response(
        status=200,
        content_type='application/json',
    )
    response.write(json.dumps({'pk': pk}))
    response.finish()
# Second Done
# First Done
```

# Why Does That Work?

# Why Does That Work?

- Still *one* process/thread

# Why Does That Work?

- Still *one* process/thread
- While one request waits for data, handle other requests

# An Analogy: A Restaurant

# Story Time

# Story Time

- Converted the most used APIs out of our app to use async IO

# Story Time

- Converted the most used APIs of our app to use async IO
  - Heavy use of threads for legacy code
  - Tornado event loop
  - Still Python 2.7



# Story Time

- Converted the most used APIs of our app to use async IO
  - Heavy use of threads for legacy code
  - Tornado event loop
  - Still Python 2.7
- Reduced number of servers by **25%**

# Story Time

- Converted the most used APIs out of our app to use async IO
  - Heavy use of threads for legacy code
  - Tornado event loop
  - Still Python 2.7
- Reduced number of servers by **25%**
  - Server bill reduced by thousands of dollars per month
  - Business logic stayed the same
  - Only moved IO code to non-blocking or threads

# Why You Might Want To Go Async

~~Why You Might Want To Go Async~~

~~Why You Might Want To Go Async~~

Why You Might Want To Go Async Now

# Why You Might Want To Go Async Now

# Why You Might Want To Go Async Now

- Python 3.6
  - `async`
  - `await`
  - `async` iterators
  - `async` generators
  - `async` context managers

# Why You Might Want To Go Async Now

- Python 3.6
  - `async`
  - `await`
  - `async` iterators
  - `async` generators
  - `async` context managers
- Async Revolution
  - `uvloop`
  - `aiopg`
  - `sanic`
  - `aio*`



# Story Time (cont)

# Story Time (cont)

- Upgraded App To Python 3.6

# Story Time (cont)

- Upgraded App To Python 3.6
  - Replaced Tornado Main Loop with AsyncIO Main Loop

# Story Time (cont)

- Upgraded App To Python 3.6
  - Replaced Tornado Main Loop with AsyncIO Main Loop
  - Removed ~20 dependencies

# Story Time (cont)

- Upgraded App To Python 3.6
  - Replaced Tornado Main Loop with AsyncIO Main Loop
  - Removed ~20 dependencies
  - Rewrote 2 dependencies (Google API, APNS)

# Story Time (cont)

- Upgraded App To Python 3.6
  - Replaced Tornado Main Loop with AsyncIO Main Loop
  - Removed ~20 dependencies
  - Rewrote 2 dependencies (Google API, APNS)
  - Used *some* asyncio libraries

# Story Time (cont)

- Upgraded App To Python 3.6
  - Replaced Tornado Main Loop with AsyncIO Main Loop
  - Removed ~20 dependencies
  - Rewrote 2 dependencies (Google API, APNS)
  - Used *some* asyncio libraries
- Result

# Story Time (cont)

- Upgraded App To Python 3.6
  - Replaced Tornado Main Loop with AsyncIO Main Loop
  - Removed ~20 dependencies
  - Rewrote 2 dependencies (Google API, APNS)
  - Used *some* asyncio libraries
- Result
  - **3x** speedup in request handling



# Story Time (cont)

- Upgraded App To Python 3.6
  - Replaced Tornado Main Loop with AsyncIO Main Loop
  - Removed ~20 dependencies
  - Rewrote 2 dependencies (Google API, APNS)
  - Used *some* asyncio libraries
- Result
  - **3x** speedup in request handling
  - Reduced number of servers by **another 30%**

# To Recap

# To Recap

- Use Async To Save Money

# To Recap

- Use Async To Save Money
- Use Async **Now** Because It's Ready

# To Recap

- Use Async To Save Money
- Use Async **Now** Because It's Ready
- Use Python 3.6

# Questions?

## Thank You!

[@ojiidotch](#)

[github.com/ojii](https://github.com/ojii)